5.2 Clock Interrupt Handling

- Clock interrupt is the 2nd to the power-failure interrupt.
- Tasks:
  - Rearms the hardware clock if necessary
  - Update CPU usage statistics
  - Performs scheduler-related functions
  - Sends a SIGXCPU signal to the current process
  - Updates the time-of-day and other related clocks.
  - Handles callouts
  - Wakes up system processes
  - Handles alarms

7

5.3 Scheduler Goals

- The scheduler must ensure that the system delivers acceptable performance to each application.
- Different applications:
  - Interactive: 50-150ms
  - Batch: scientific computation
  - Real-time: time-critical

11

5.4 Traditional UNIX Scheduling

- To improve response times of interactive users, while ensuring that low-priority, background jobs do not starve.
- Priority-based:
  - User-process is preempted
  - Kernel is strictly non-preemptive

12

Priority

- Kernel:0-49, user: 50-127
- proc fields:
  - p_pri: Current scheduling priority
  - p_usrpri: User mode priority
  - p_cpu: Measure of recent CPU usage
  - p_nice: User-controllable nice factor
- Kernel:
  - Sleeping priority

13

User mode priority

- Depends on two factors:
  - Nice: 0-39
  - CPU usage
- Time-sharing: equal opportunity
- decay factor: for SVR3 it is 1/2, for 4.3BSD:
  - decay = (2*load_average)/(2*load_average+1)
  - p_cpu = p_cpu* decay
- p_usrpri = PUSER + (p_cpu/4) +(2*p_nice)

14

Example : PUSER = 50

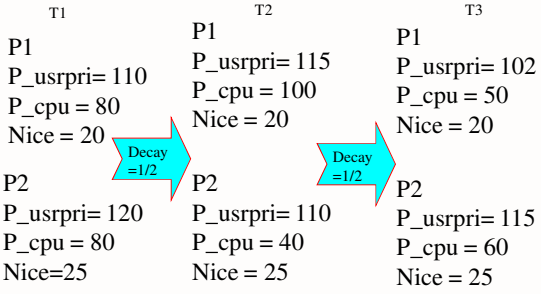

15

Scheduler Implementation

- 32 run queues: doubly linked list of proc structures for runnable processes.
- whichqs: bitmask for each queue, “1” means that there is a runnable process
- swtch(): context switch by p_addr
  - Saving part of u area (pcb)
  - Loading the saved context.
- VAX ffs, ffc, INSQHI, REMQHI, LDPCTX, SVPCTX : special instructions for context switch etc.

16

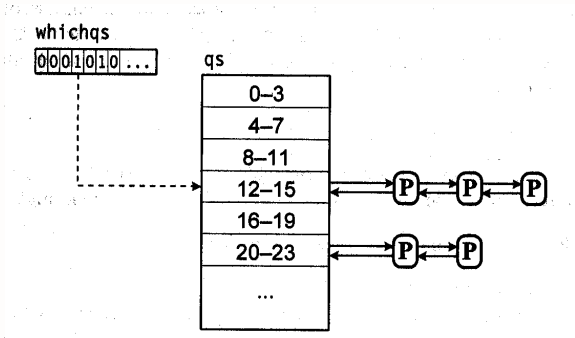

Figure 5-2. BSD scheduler data structures.

17